

Homework 3

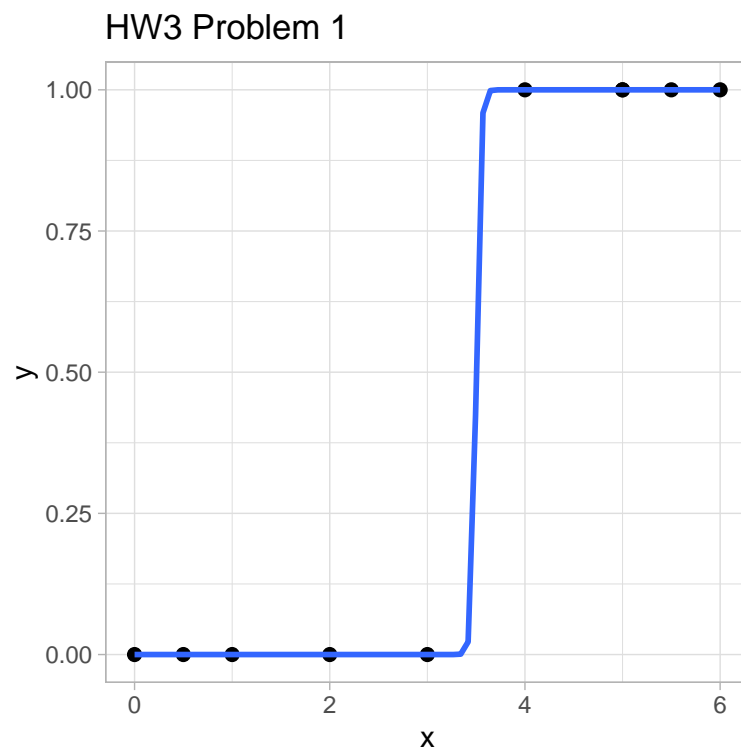
Ericka Smith

10/14/2020

Problem 1

```
dat <- data.frame(  
  x = c(0, 0.5, 1, 2, 3, 4, 5, 5, 5.5, 6),  
  y = c(0, 0, 0, 0, 0, 1, 1, 1, 1, 1))  
  
ggplot(data = dat, aes(x = x, y=y))+  
  geom_point(size = 2)+  
  geom_smooth(method = "glm", method.args = list(family=binomial), se=F)+  
  theme_light()+  
  ggtitle("HW3 Problem 1")
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```
summary(glm(y~x, family = binomial, data=dat))
```

```
##
## Call:
## glm(formula = y ~ x, family = binomial, data = dat)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.575e-05  -2.110e-08   0.000e+00   2.110e-08   1.607e-05
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -159.55  270391.50  -0.001      1
## x              45.58   76267.22   0.001      1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1.3863e+01  on 9  degrees of freedom
## Residual deviance: 5.0616e-10  on 8  degrees of freedom
## AIC: 4
##
## Number of Fisher Scoring iterations: 25
```

For both the intercept and the coefficient for x we get p -values equal to 1, so neither estimate is significant and we can't interpret their values in a very meaningful way. That being said, in a logistic regression output such as this the intercept signifies the log odds of being in the $p(x) = 0$ group for y , and the coefficient of x signifies the change in the log odds of being in the $p(x) = 1$ group as compared to the $p(x) = 0$ group.

Problem 2

Part A They are equivalent as parameter points because the constant c falls away in the model, because $\alpha_i = \log(w_i)$ and $+c$ is the only difference between $\alpha_1, \dots, \alpha_8$ as compared to $\alpha_1 + c, \dots, \alpha_8 + c$. We can still estimate the weights w_1, \dots, w_8 because we have the relationship between these already.

```
X <- read_csv("p2.csv", col_names = FALSE)
```

Part B

```
## Parsed with column specification:
## cols(
##   X1 = col_double(),
##   X2 = col_double(),
##   X3 = col_double(),
##   X4 = col_double(),
##   X5 = col_double(),
##   X6 = col_double(),
##   X7 = col_double()
## )
```

```
(xmat <- as.matrix(X))
```

```
##      X1 X2 X3 X4 X5 X6 X7
## [1,]  1 -1  0  0  0  0  0
## [2,]  1  0 -1  0  0  0  0
## [3,]  1  0  0 -1  0  0  0
## [4,]  1  0  0  0 -1  0  0
## [5,]  1  0  0  0  0 -1  0
## [6,]  1  0  0  0  0  0 -1
## [7,]  1  0  0  0  0  0  0
## [8,]  0  1 -1  0  0  0  0
## [9,]  0  1  0 -1  0  0  0
## [10,] 0  1  0  0 -1  0  0
## [11,] 0  1  0  0  0 -1  0
## [12,] 0  1  0  0  0  0 -1
## [13,] 0  1  0  0  0  0  0
## [14,] 0  0  1 -1  0  0  0
## [15,] 0  0  1  0 -1  0  0
## [16,] 0  0  1  0  0 -1  0
## [17,] 0  0  1  0  0  0 -1
## [18,] 0  0  1  0  0  0  0
## [19,] 0  0  0  1 -1  0  0
## [20,] 0  0  0  1  0 -1  0
## [21,] 0  0  0  1  0  0 -1
## [22,] 0  0  0  1  0  0  0
## [23,] 0  0  0  0  1 -1  0
## [24,] 0  0  0  0  1  0 -1
## [25,] 0  0  0  0  1  0  0
## [26,] 0  0  0  0  0  1 -1
## [27,] 0  0  0  0  0  1  0
## [28,] 0  0  0  0  0  0  1
```

```
y <- c(5,3,7,6,6,1/3,1/4,1/3,5,3,3,1/5,1/7,6,3,4,6,1/5,1/3,1/4,1/7,1/8,1/2,1/5,1/6,1/5,1/6,1/2)
logy <- map_dbl(y, ~log(.x))
model0 <- lm(y~xmat)
summary(model0)
```

```
##
## Call:
## lm(formula = y ~ xmat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7216 -0.7734  0.0876  0.7992  3.1946
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.4300     0.5442   4.465 0.000237 ***
## xmatX1         -0.6107     1.1938  -0.512 0.614550
## xmatX2         -2.6166     1.0884  -2.404 0.026020 *
## xmatX3         -0.8603     0.9904  -0.869 0.395366
## xmatX4         -4.3797     0.9024  -4.853 9.64e-05 ***
## xmatX5         -3.0619     0.8275  -3.700 0.001417 **
```

```
## xmatX6      -2.6940      0.7696  -3.501  0.002252 **
## xmatX7      -1.2356      0.7326  -1.687  0.107222
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.44 on 20 degrees of freedom
## Multiple R-squared:  0.7459, Adjusted R-squared:  0.6569
## F-statistic: 8.386 on 7 and 20 DF,  p-value: 8.272e-05
```

Here are the estimates I obtain:

```
 $\alpha_1 - \alpha_2 = -0.6107$ 
 $\alpha_1 - \alpha_3 = -2.6166$ 
 $\alpha_1 - \alpha_4 = -0.8603$ 
 $\alpha_1 - \alpha_5 = -4.3797$ 
 $\alpha_1 - \alpha_6 = -3.0619$ 
 $\alpha_1 - \alpha_7 = -2.6940$ 
 $\alpha_1 - \alpha_8 = -1.2356$ 
```

```
exp_a <- unname(exp(model0$coefficients))
w <- rbind(diag(-1,7,7), rep(1,7))
(w <- cbind(w, as.vector(exp_a)))
```

Part C

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7]      [,8]
## [1,]  -1    0    0    0    0    0    0 11.35907526
## [2,]   0   -1    0    0    0    0    0  0.54296290
## [3,]   0    0   -1    0    0    0    0  0.07305060
## [4,]   0    0    0   -1    0    0    0  0.42303976
## [5,]   0    0    0    0   -1    0    0  0.01252887
## [6,]   0    0    0    0    0   -1    0  0.04679648
## [7,]   0    0    0    0    0    0   -1  0.06760818
## [8,]   1    1    1    1    1    1    1  0.29065181
```

```
(z <- c(rep(0,7), 1))
```

```
## [1] 0 0 0 0 0 0 0 1
```

```
solve(w,z)
```

```
## [1] 0.8863396439 0.0423669649 0.0057000807 0.0330094571 0.0009776181
## [6] 0.0036514923 0.0052754124 0.0780292078
```

The estimates of the weights are the following:

```
 $w_1 = 0.8863396439$ 
```

$$w_2 = 0.0423669649$$

$$w_3 = 0.0057000807$$

$$w_4 = 0.0330094571$$

$$w_5 = 0.0009776181$$


$$w_6 = 0.0036514923$$

$$w_7 = 0.0052754124$$


$$w_8 = 0.0780292078$$


Part D Since only certain α estimates were significant that calls into question the hypothesis that all 8 criteria are equally important. The varying magnitudes of the weights also do. For example, w_5 is a LOT smaller than the other weights, which leads me to believe that yard space (5) is not actually as important as some of the other criteria.


```
plot(model0)
```



Part E







Based on the plots the model assumptions are not quite met. In the Residuals vs. Fitted plot you can see that there is nonconstant variance and on the Q-Q plot there is a bit of tailing off at the upper right hand corner of the plot.

Problem 3

```
# dat <- data.frame(
#   insecticide = c(2, 2.64, 3.48, 4.59, 6.06, 8),
#   ddt = c(3/50, 5/49, 19/47, 19/38, 24/49, 35/50),
#   bhc = c(2/50, 14/49, 20/50, 27/50, 41/50, 40/50),
#   both = c(28/50, 37/50, 46/50, 48/50, 48/50, 50/50))
df <- data.frame(
  dose <- c(2, 2.64, 3.48, 4.59, 6.06, 8, 2, 2.64, 3.48, 4.59, 6.06, 8, 2, 2.64, 3.48, 4.59, 6.06, 8),
  kr <- c(3/50, 5/49, 19/47, 19/38, 24/49, 35/50, 2/50, 14/49, 20/50, 27/50, 41/50, 40/50, 28/50, 37/50, 46/50, 48/50, 48/50, 50/50),
  insecticide <- c("DDT", "DDT", "DDT", "DDT", "DDT", "DDT", "BHC", "BHC", "BHC", "BHC", "BHC", "BHC", "BHC", "BHC", "BHC", "BHC", "BHC", "BHC"))

ggplot(data = df)+
  geom_smooth(aes(x = dose, y=kr, color=insecticide), method = "glm", method.args = list(family=binomial)) +
  geom_point(aes(x = dose, y=kr, color=insecticide))+
  theme_light()+
  ggtitle("Toxicity of Insecticides")+
  xlab(as.expression(bquote("Insecticide Dose (mg/10" ~ cm^2 ~ ")")))+
  ylab("Kill Rate")
```

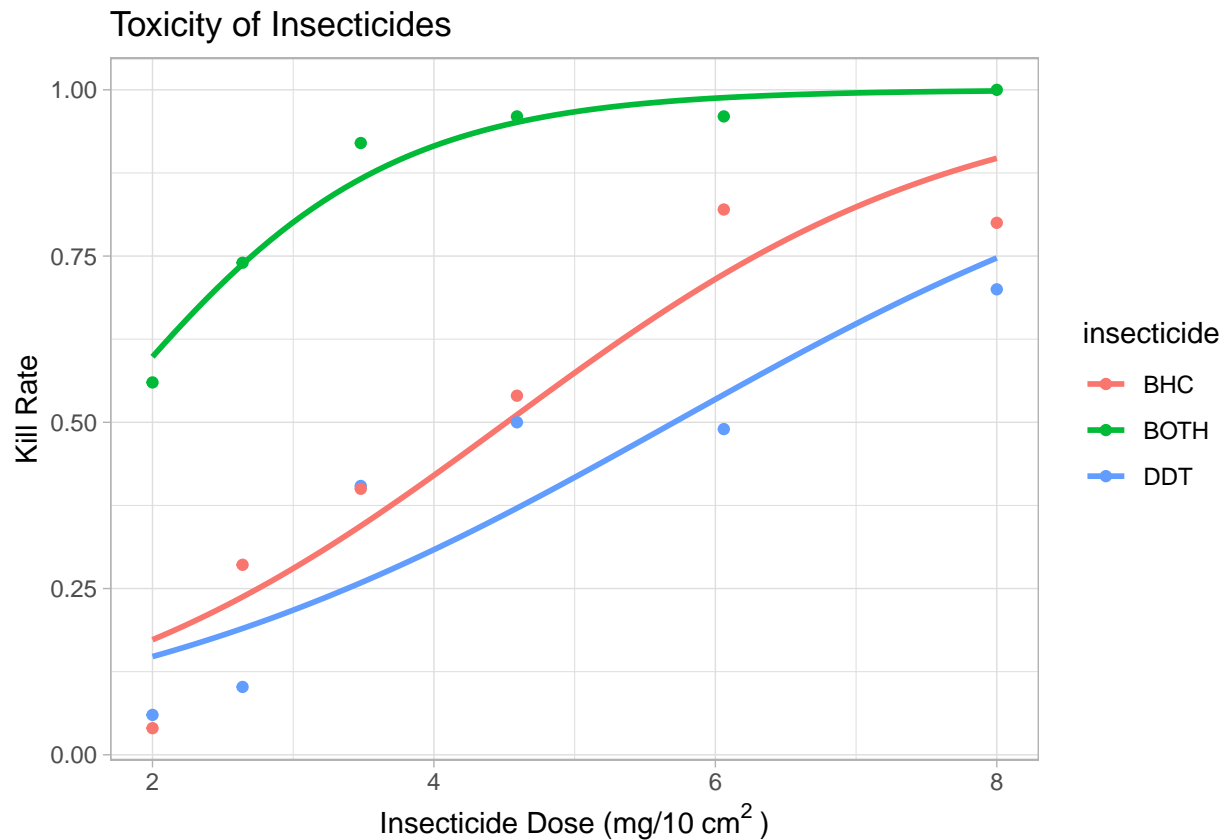

Parts A and B

```
## `geom_smooth()` using formula 'y ~ x'
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```



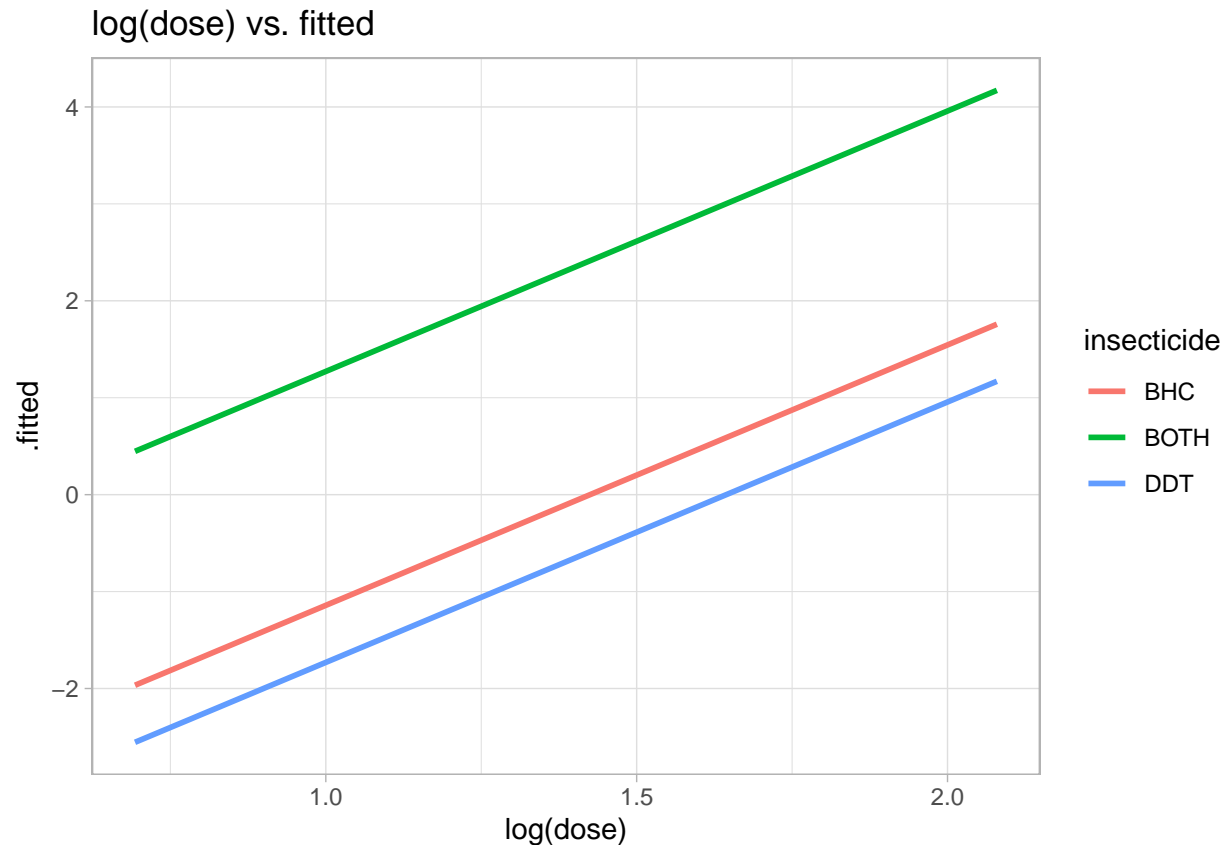
```
mod <- glm(kr ~ log(dose) + insecticide, family="binomial")
```

Part C

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```

```
ggplot(mod)+  
  geom_smooth(aes(log(dose), .fitted, color = insecticide), se = F)+  
  theme_light()+  
  ggtitle("log(dose) vs. fitted")
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Looking at the plot of log(dose) vs. fitted values we see that these lines are parallel. This gives evidence to the hypothesis of parallelism.

Part D In the formula chem+ldose there is an intercept term and in the formula chem+ldose-1 there is not. Since this is the only difference the covariance matrices will be the same.

```
a = mod$coefficients[1]
b = mod$coefficients[2]
V=vcov(mod)
z = 1.645;
k = 1.645^2*V[2,2]/b^2
(tau = polyroot(c(a^2 - V[1,1]*z*z, 2*a*b - 2*V[1,2]*z*z, b^2 - V[2,2]*z*z)))
```

Part E

```
## [1] -0.0048049+0i 3.0788096+0i
```

The 90% Confidence Interval that I obtain is (-0.0048049, 3.0788096)

```
mod_probit <- glm(kr ~ log(dose) + insecticide, family=binomial(link = "probit"))
```

Part F

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```

```
mod_cloglog <- glm(kr ~ log(dose) + insecticide, family=binomial(link = "cloglog"))
```


```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```

```
mod_loglog <- glm((1-kr) ~ log(dose) + insecticide, family=binomial(link="cloglog"))
```

```
## Warning in eval(family$initialize): non-integer #successes in a binomial glm!
```


```
ggplot(mod_probit)+  
  geom_smooth(aes(log(dose), .fitted, color = insecticide), se = F)+  
  theme_light()+  
  ggtitle("probit model")
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




```
ggplot(mod_cloglog)+
  geom_smooth(aes(log(dose), .fitted, color = insecticide), se = F)+
  theme_light()+
  ggtitle("c-log-log model")
```

`geom_smooth()` using method = 'loess' and formula 'y ~ x'



```
ggplot(mod_loglog)+
  geom_smooth(aes(log(dose), .fitted, color = insecticide), se = F)+
  theme_light()+
  ggtitle("log-log model")
```

`geom_smooth()` using method = 'loess' and formula 'y ~ x'



```
a = mod_cloglog$coefficients[1]
b = mod_cloglog$coefficients[2]
V=vcov(mod_cloglog)
z = 1.645;
k = 1.645^2*V[2,2]/b^2
(tau = polyroot(c(a^2 - V[1,1]*z*z, 2*a*b - 2*V[1,2]*z*z, b^2 - V[2,2]*z*z)))
```

```
## [1] 0.532220-0i 4.451669+0i
```

It really doesn't seem to me that any of them gives an appreciably better fit. For my 90% confidence interval for the c-log-log model I get (0.532220,4.451669).

```
predict(mod, type = "response")[18]
```

Part G

```
##          18
## 0.9848035
```

```
exp(predict(mod)[18])
```

```
##          18
## 64.80447
```

The closest we can reliably get to 0.99 is 0.98480346, and this occurs at a dose of $64.8mg/10cm^2$

Part H It seems pretty clear to me based upon the prior plots and models that both insecticides together are the most effective. DDT appears to be less effective than gamma-BHC but the two are fairly similar overall.